

```
from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

↗ Mounted at /content/drive

```
!pip install torch torchvision efficientnet_pytorch matplotlib
```

↗ Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: torchvision in /usr/local/lib/python3.10/dis
Collecting efficientnet_pytorch
 Downloading efficientnet_pytorch-0.7.1.tar.gz (21 kB)
 Preparing metadata (setup.py) ... done
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/p
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/pyth
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.1
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/di
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.1
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/pytho
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-p
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3
Building wheels for collected packages: efficientnet_pytorch
 Building wheel for efficientnet_pytorch (setup.py) ... done
 Created wheel for efficientnet_pytorch: filename=efficientnet_pytorch-0.7
 Stored in directory: /root/.cache/pip/wheels/03/3f/e9/911b1bc46869644912b
Successfully built efficientnet_pytorch
Installing collected packages: efficientnet_pytorch
Successfully installed efficientnet_pytorch-0.7.1

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, random_split
from torchvision import datasets, transforms
from efficientnet_pytorch import EfficientNet
import matplotlib.pyplot as plt
import numpy as np
from collections import Counter
import os
import warnings as w
w.filterwarnings('ignore')
```

```
from google.colab import drive
drive.mount('/content/drive')
```

↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call

```
# Image transformations for augmentation, especially for classes with fewer samples
train_transforms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.RandomRotation(15),
    transforms.RandomHorizontalFlip(),
    transforms.RandomVerticalFlip(),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
])
```

```
test_transforms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
])
```

```
dataset = datasets.ImageFolder("/content/drive/MyDrive/Food Classification data")
```

```
# Split the dataset into training and testing sets (e.g., 80% train, 20% test)
train_size = int(0.8 * len(dataset))
test_size = len(dataset) - train_size
train_data, test_data = random_split(dataset, [train_size, test_size])
```

```
# Apply test transforms to the test dataset
test_data.dataset.transform = test_transforms
```

```

train_labels = [dataset.targets[i] for i in train_data.indices]
class_counts = Counter(train_labels)
print("Class counts in training data:", class_counts)

# Oversampling or data augmentation strategy based on class imbalance
class_weights = 1. / np.array([class_counts[i] for i in range(len(class_counts))])
sample_weights = [class_weights[label] for label in train_labels]

# Weighted sampler for handling imbalanced classes
weighted_sampler = torch.utils.data.WeightedRandomSampler(sample_weights, len(s

⇒ Class counts in training data: Counter({7: 815, 2: 812, 8: 797, 27: 796, 21

batch_size = 32

train_loader = DataLoader(train_data, batch_size=batch_size, sampler=weighted_s
test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=False)

model = EfficientNet.from_pretrained('efficientnet-b0')

# Modify the final layer to match the number of food classes
num_classes = 28
model._fc = nn.Linear(model._fc.in_features, num_classes)

# Move the model to GPU if available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = model.to(device)

⇒ Downloading: "https://github.com/lukemelas/EfficientNet-PyTorch/releases/download/v0.7.1/efficientnet-b0-380332dd.pth"
100%|██████████| 20.4M/20.4M [00:00<00:00, 384MB/s]
Loaded pretrained weights for efficientnet-b0

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

def train_model(model, criterion, optimizer, train_loader, test_loader, num_epochs):
    train_loss_history = []
    test_loss_history = []
    train_acc_history = []
    test_acc_history = []

    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        correct = 0

```

```
total = 0

for images, labels in train_loader:
    images, labels = images.to(device), labels.to(device)

    optimizer.zero_grad()
    outputs = model(images)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()

    running_loss += loss.item()
    _, predicted = torch.max(outputs, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

train_loss_history.append(running_loss / len(train_loader))
train_acc_history.append(100 * correct / total)

# Validation
model.eval()
val_loss = 0.0
val_correct = 0
val_total = 0
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        loss = criterion(outputs, labels)
        val_loss += loss.item()

        _, predicted = torch.max(outputs, 1)
        val_total += labels.size(0)
        val_correct += (predicted == labels).sum().item()

test_loss_history.append(val_loss / len(test_loader))
test_acc_history.append(100 * val_correct / val_total)

print(f'Epoch [{epoch+1}/{num_epochs}], '
      f'Train Loss: {running_loss/len(train_loader):.4f}, '
      f'Train Acc: {100 * correct / total:.2f}%, '
      f'Val Loss: {val_loss/len(test_loader):.4f}, '
      f'Val Acc: {100 * val_correct / val_total:.2f}%')

return train_loss_history, test_loss_history, train_acc_history, test_acc_h
```

```
num_epochs = 10
train_loss, test_loss, train_acc, test_acc = train_model(model, criterion, opti

def plot_metrics(train_loss, test_loss, train_acc, test_acc):
    epochs = range(1, len(train_loss) + 1)

    plt.figure(figsize=(12, 4))

    # Plot loss
    plt.subplot(1, 2, 1)
    plt.plot(epochs, train_loss, 'b', label='Train Loss')
    plt.plot(epochs, test_loss, 'r', label='Test Loss')
    plt.title('Training and Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()

    # Plot accuracy
    plt.subplot(1, 2, 2)
    plt.plot(epochs, train_acc, 'b', label='Train Accuracy')
    plt.plot(epochs, test_acc, 'r', label='Test Accuracy')
    plt.title('Training and Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy (%)')
    plt.legend()

    plt.show()

# Cell 11: Call the plotting function
plot_metrics(train_loss, test_loss, train_acc, test_acc)
```



```
def continue_training_with_early_stopping(model, criterion, optimizer, train_lc
    best_acc = 0.0 # To track the best validation accuracy
    for epoch in range(start_epoch, start_epoch + num_more_epochs):
        model.train()
        running_loss = 0.0
        correct = 0
```

```

total = 0

for images, labels in train_loader:
    images, labels = images.to(device), labels.to(device)

    optimizer.zero_grad()
    outputs = model(images)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()

    running_loss += loss.item()
    _, predicted = torch.max(outputs, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

# Validation phase
model.eval()
val_loss = 0.0
val_correct = 0
val_total = 0
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        loss = criterion(outputs, labels)
        val_loss += loss.item()

        _, predicted = torch.max(outputs, 1)
        val_total += labels.size(0)
        val_correct += (predicted == labels).sum().item()

val_acc = 100 * val_correct / val_total

print(f'Epoch [{epoch+1}/{start_epoch + num_more_epochs}], '
      f'Train Acc: {100 * correct / total:.2f}%, '
      f'Val Loss: {val_loss/len(test_loader):.4f}, '
      f'Val Acc: {val_acc:.2f}%')

# Early stopping if validation accuracy reaches or exceeds the threshold
if val_acc >= early_stop_acc:
    print(f'Early stopping triggered at epoch {epoch+1} with validation accuracy {val_acc:.2f}%')
    break

# Continue training with early stopping
start_epoch = 10 # Continue from the 6th epoch
num_more_epochs = 10 # Number of additional epochs to train
continue_training_with_early_stopping(model, criterion, optimizer, train_loader,

```

```
↩ Epoch [11/20], Train Acc: 96.58%, Val Loss: 0.7156, Val Acc: 82.90%  
Epoch [12/20], Train Acc: 96.93%, Val Loss: 0.6305, Val Acc: 84.06%  
Epoch [13/20], Train Acc: 97.02%, Val Loss: 0.7211, Val Acc: 82.32%  
Epoch [14/20], Train Acc: 97.09%, Val Loss: 0.7753, Val Acc: 80.04%  
Epoch [15/20], Train Acc: 97.32%, Val Loss: 0.8426, Val Acc: 82.07%  
Epoch [16/20], Train Acc: 98.22%, Val Loss: 0.5235, Val Acc: 86.00%  
Early stopping triggered at epoch 16 with validation accuracy: 86.00%
```

```
# Cell 13: Save the trained model
```

```
torch.save(model.state_dict(), 'food_classification_model.pth')  
print("Model saved as 'food_classification_model.pth'")
```

```
↩ Model saved as 'food_classification_model.pth'
```

```
import os  
import pickle
```

```
# Define the path to the dataset
```

```
dataset_path = "/content/drive/MyDrive/Food Classification dataset1/"
```

```
# Get class names from folder names
```

```
class_names = sorted(os.listdir(dataset_path))
```

```
# Create a dictionary mapping class indices to class names
```

```
class_mapping = {idx: class_name for idx, class_name in enumerate(class_names)}
```

```
# Save the class mapping as a pickle file
```

```
with open('class_mapping.pkl', 'wb') as f:  
    pickle.dump(class_mapping, f)
```

class_mapping

```
↔ {0: 'Donut',  
  1: 'Hot Dog',  
  2: 'apple_pie',  
  3: 'burger',  
  4: 'butter_naam',  
  5: 'chai',  
  6: 'chapati',  
  7: 'cheesecake',  
  8: 'chicken_curry',  
  9: 'chole_bhature',  
 10: 'dal_makhani',  
 11: 'dhokla',  
 12: 'fried_rice',  
 13: 'ice_cream',  
 14: 'idli',  
 15: 'jalebi',  
 16: 'kaathi_rolls',  
 17: 'kadai_paneer',  
 18: 'kulfi',  
 19: 'masala_dosa',  
 20: 'momos',  
 21: 'omelette',  
 22: 'paani_puri',  
 23: 'pakode',  
 24: 'pav_bhaji',  
 25: 'pizza',  
 26: 'samosa',  
 27: 'sushi'}
```